

1 Machine Learning Formulae

Note: Lecture notes and slides available here:

<https://vkosuri.github.io/CourseraMachineLearning/>

1.1 Notation

1.1.1 Linear and logistic regression

Feature:

$$x_j \quad (1)$$

Single data point in a feature:

$$x_j^{(i)} \quad (2)$$

Feature vector:

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix} \quad (3)$$

Matrix of training examples, stored row-wise:

$$X = \begin{bmatrix} x_0 & x_1 & x_2 & \cdots & x_n \end{bmatrix} = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} & x_2^{(1)} & \cdots & x_n^{(1)} \\ x_0^{(2)} & x_1^{(2)} & x_2^{(2)} & \cdots & x_n^{(2)} \\ x_0^{(3)} & x_1^{(3)} & x_2^{(3)} & \cdots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_0^{(m)} & x_1^{(m)} & x_2^{(m)} & \cdots & x_n^{(m)} \end{bmatrix} \quad (4)$$

Theta:

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \quad (5)$$

Feature scaling:

$$x_j := \frac{x_j}{s_j} \quad \text{where} \quad s_j = \max(x_j) - \min(x_j) \quad (6)$$

Mean normalisation:

$$x_j := x_j - \mu_j \quad \text{where} \quad \mu_j = \frac{\sum_{i=1}^m x_j^{(i)}}{m} \quad (7)$$

Feature normalisation (feature scaling *and* mean normalisation):

$$x_j := \frac{x_j - \mu_j}{s_j} \quad (8)$$

1.1.2 Neural networks

General structure:

$$\begin{bmatrix} x_0 \\ x_1 \\ x_3 \\ \vdots \\ x_n \end{bmatrix} \rightarrow \begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ a_2^{(1)} \\ \vdots \\ a_n^{(1)} \end{bmatrix} \rightarrow \begin{bmatrix} a_0^{(2)} \\ a_1^{(2)} \\ a_2^{(2)} \\ \vdots \\ a_n^{(2)} \end{bmatrix} \quad (9)$$

1.2 Linear regression

1.2.1 Basic equations

Hypothesis function:

$$h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n \quad (10)$$

Cost function:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}^{(i)} - y^{(i)})^2 \quad (11)$$

$$J(\theta) = \frac{1}{2m} (X\theta - y)^T (X\theta - y) \quad (12)$$

Gradient descent:

$$\theta_j := \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j} \quad (13)$$

$$\theta := \theta - \alpha \nabla J(\theta) \quad (14)$$

Gradient descent (update rule):

$$\begin{aligned} & \text{repeat until convergence : } \{ \\ & \quad \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j \\ & \} \end{aligned} \quad (15)$$

$$\begin{aligned} & \text{repeat until convergence : } \{ \\ & \quad \theta := \theta - \frac{\alpha}{m} X^T (X\theta - y) \\ & \} \end{aligned} \quad (16)$$

Normal equation:

$$\theta = (X^T X)^{-1} X^T y \quad (17)$$

1.2.2 Polynomial regression

Hypothesis function:

$$h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n \quad (18)$$

But where some of the higher-order features are more complex functions of lower order features e.g.:

$$x_3 = x_1 \sqrt{x_2} \quad (19)$$

1.2.3 Equations with regularisation

Cost function:

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \quad (20)$$

$$J(\theta) = \frac{1}{2m} [(X\theta - y)^T (X\theta - y) + \lambda \theta^T \theta] \quad (21)$$

Gradient descent (update rule):

$$\begin{aligned} & \text{repeat until convergence : } \{ \\ & \quad \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_0 \\ & \quad \theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j \right) + \frac{\lambda}{m} \theta_j \right] \\ & \quad \} \quad (22) \end{aligned}$$

Note: By convention, the θ_0, x_0 update term is not regularised.

The update rule can also be re-arranged to give the original update rule, plus an additional term at the front:

$$\theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j \quad (23)$$

Normal equation:

$$\theta = (X^T X + \lambda \cdot L)^{-1} X^T y \quad (24)$$

Where:

$$L = \begin{bmatrix} 0 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix} \quad (25)$$

1.3 Logistic regression

1.3.1 Basic equations

Hypothesis function:

$$h_{\theta}(x) = g(\theta^T x) \quad (26)$$

$$h_{\theta}(x) = g(X\theta) \quad (27)$$

Where:

$$g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}} \quad (28)$$

Hypothesis function interpretation:

$$h_{\theta}(x) = P(y = 1|x; \theta) = 1 - P(y = 0|x; \theta) \quad (29)$$

Decision boundary:

$$h_{\theta}(x) \geq 0.5 \rightarrow y = 1 \quad (30)$$

$$h_{\theta}(x) < 0.5 \rightarrow y = 0 \quad (31)$$

Decision boundary (2):

$$\theta^T x \geq 0 \rightarrow y = 1 \quad (32)$$

$$\theta^T x < 0 \rightarrow y = 0 \quad (33)$$

Cost function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] \quad (34)$$

$$J(\theta) = \frac{1}{m} [-y^T \log(h) - (1 - y)^T \log(1 - h)] \quad (35)$$

Gradient descent (update rule):

repeat until convergence : {

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j \quad \} \quad (36)$$

$$\begin{aligned}
& \text{repeat until convergence : } \{ \\
& \quad \theta := \theta - \frac{\alpha}{m} X^T (g(X\theta) - y) \\
& \} \quad (37)
\end{aligned}$$

Multi-class classification:

$$\begin{aligned}
& y \in \{0, 1, 2, \dots, n\} \\
& h_{\theta}^{(0)}(x) = P(y = 0|x; \theta) \\
& h_{\theta}^{(1)}(x) = P(y = 1|x; \theta) \\
& h_{\theta}^{(2)}(x) = P(y = 2|x; \theta) \\
& \vdots \\
& h_{\theta}^{(n)}(x) = P(y = n|x; \theta) \\
& \text{prediction} = \max(h_{\theta}^{(i)}(x))
\end{aligned}$$

1.3.2 Equations with regularisation

Cost function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \quad (38)$$

$$J(\theta) = \frac{1}{m} [-y^T \log(h) - (1 - y)^T \log(1 - h)] + \frac{\lambda}{2m} \theta^T \theta \quad (39)$$

Gradient descent (update rule):

$$\begin{aligned}
& \text{repeat until convergence : } \{ \\
& \quad \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_0 \\
& \quad \theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j \right) + \frac{\lambda}{m} \theta_j \right] \\
& \} \quad (40)
\end{aligned}$$

Note: This is the same as the update rule for linear regression. By convention, the θ_0, x_0 update term is not regularised.

As with linear regression, the update rule can also be re-arranged to give the original update rule, plus an additional term at the front:

$$\theta_j := \theta_j(1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j \quad (41)$$

1.4 Neural networks

1.4.1 Notation

L = number of layers

l = specific layer number

s_l = number of units/nodes (not including bias unit) in layer l

$a_j^{(l)}$ = unit/node j in layer l

$a_j^{(t)(l)}$ = training example in unit/node j in layer l

Θ^l = matrix of weights for moving between layer l and $l + 1$

$\Theta_{i,j}^l$ = row i , column j in matrix of weights l

$a^{(1)} = X$ = input layer

$a^{(L)} = h_{\Theta}(x)$ = output layer

K = number of classes in the output layer = number of classes in y (including 0)

1.4.2 Matrix dimensions

The input layer $a^{(1)}$ is the same as the matrix of training examples X :

$$a^{(1)} = X = \left[\begin{array}{c} \leftarrow n+1 \rightarrow \\ \uparrow \\ m \\ \downarrow \end{array} \right] = \left[\begin{array}{cccccc} x_0 & x_1 & x_2 & \cdots & x_n \end{array} \right] \quad (42)$$

Where:

$$x_0 = [1 \quad 1 \quad 1 \quad \cdots \quad 1] \quad (43)$$

The hidden layers have $(s_l + 1)$ columns (the '+1' is the bias unit), each with m training examples:

$$a^{(l)} = \left[\begin{array}{c} \leftarrow (s_l + 1) \rightarrow \\ \uparrow \\ m \\ \downarrow \end{array} \right] = \left[\begin{array}{cccccc} a_0^{(l)} & a_1^{(l)} & a_2^{(l)} & \cdots & a_n^{(l)} \end{array} \right] \quad (44)$$

Where:

$$a_0^{(l)} = [1 \quad 1 \quad 1 \quad \cdots \quad 1] \quad (45)$$

The output layer $a^{(L)}$ is the model hypothesis $h_{\Theta}(x)$.

For a single class neural network, this is a vector whose values correspond to the hypothesis value for each entry in the training data:

$$a^{(L)} = h_{\Theta}(x) = \begin{bmatrix} h_{(\Theta)}(x^{(1)}) \\ h_{(\Theta)}(x^{(2)}) \\ h_{(\Theta)}(x^{(3)}) \\ \vdots \\ h_{(\Theta)}(x^{(m)}) \end{bmatrix} \quad (46)$$

For a multiclass neural network, this is a matrix of values whose columns correspond to the different classes, and whose rows correspond to each entry in the training data:

$$a^{(L)} = h_{\Theta}(x) = \begin{bmatrix} h_{\Theta}(x)_1 \\ h_{\Theta}(x)_2 \\ h_{\Theta}(x)_3 \\ \vdots \\ h_{\Theta}(x)_K \end{bmatrix}$$

$$= \begin{bmatrix} \leftarrow K \rightarrow \\ \uparrow \\ m \\ \downarrow \end{bmatrix} = \begin{bmatrix} h_{\Theta}(x^{(t)})_1 & h_{\Theta}(x^{(t)})_2 & h_{\Theta}(x^{(t)})_3 & \cdots & h_{\Theta}(x^{(t)})_K \end{bmatrix} \quad (47)$$

Remember, the hypothesis values correspond to the probability of training example t being in class K :

$$h_{\Theta}(x^{(t)})_k = P(x^{(t)} = k|x; \Theta) \quad (48)$$

The matrix of weights $\Theta_{i,j}^l$ has $s_l + 1$ rows (the +1 from the bias unit) and

(s_{l+1}) columns:

$$\Theta_{i,j}^l = \left[\begin{array}{c} \leftarrow s_l + 1 \rightarrow \\ \uparrow \\ s_{l+1} \\ \downarrow \end{array} \right] \quad (49)$$

The data classifications in y are represented as a binary $m \times k$ matrix, with each row corresponding to an entry in the training data, and each column corresponding to a different class.

$$y = \left[\begin{array}{ccccc} y_0 & y_1 & y_2 & \cdots & y_K \end{array} \right] = \left[\begin{array}{c} \leftarrow K \rightarrow \\ \uparrow \\ m \\ \downarrow \end{array} \right] \quad (50)$$

And the mapping is as follows:

$$y = \begin{bmatrix} 1 \\ 5 \\ 2 \\ 8 \\ 9 \\ 0 \\ 5 \\ \vdots \\ 9 \end{bmatrix} \rightarrow \begin{bmatrix} 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix} \quad (51)$$

1.4.3 Basic representation

Single class neural network:

$$\begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \rightarrow \begin{bmatrix} a_0^{(2)} \\ a_1^{(2)} \\ a_2^{(2)} \\ \dots \end{bmatrix} \rightarrow \begin{bmatrix} a_0^{(3)} \\ a_1^{(3)} \\ a_2^{(3)} \\ \dots \end{bmatrix} \rightarrow \dots \rightarrow \begin{bmatrix} a_0^{(j)} \\ a_1^{(j)} \\ a_2^{(j)} \\ \dots \end{bmatrix} \rightarrow \dots \rightarrow h_{\Theta}(x) \quad (52)$$

Multiclass neural network:

$$\begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \rightarrow \begin{bmatrix} a_0^{(2)} \\ a_1^{(2)} \\ a_2^{(2)} \\ \dots \end{bmatrix} \rightarrow \begin{bmatrix} a_0^{(3)} \\ a_1^{(3)} \\ a_2^{(3)} \\ \dots \end{bmatrix} \rightarrow \dots \rightarrow \begin{bmatrix} a_0^{(j)} \\ a_1^{(j)} \\ a_2^{(j)} \\ \dots \end{bmatrix} \rightarrow \dots \rightarrow \begin{bmatrix} h_{\Theta}(x)_1 \\ h_{\Theta}(x)_2 \\ h_{\Theta}(x)_3 \\ \vdots \\ h_{\Theta}(x)_k \end{bmatrix} \quad (53)$$

1.4.4 Forward propagation

The role of feed forward propagation is to calculate the hypothesis $h_{\Theta}(x)$ and - if tracking - the overall cost $J(\Theta)$. The values of $h_{\Theta}(x)$ feed directly into the back propagation algorithm.

Note: Feed forward and back propagation are carried out element-wise through the training data, typically using a for-loop:

$$\begin{aligned} &\text{for } i \text{ in range(m):} \\ &\dots \\ &\dots \end{aligned} \quad (54)$$

For layer 1:

$$a^{(1)} = X \quad (55)$$

For layer 2:

$$\begin{aligned} z^{(2)} &= \Theta^{(1)T} X \\ a^{(2)} &= g(z^{(2)}) \end{aligned}$$

$$[\text{Then add bias unit as column of ones}] \quad (56)$$

For layer $l + 1$:

$$\begin{aligned} z^{(l+1)} &= \Theta^{(l)T} a^{(l)} \\ a^{(l+1)} &= g(z^{(l+1)}) \end{aligned}$$

$$[\text{Then add bias unit as column of ones}] \quad (57)$$

For output layer:

$$\begin{aligned} z^{(L)} &= \Theta^{(L-1)T} a^{(L-1)} \\ a^{(L)} &= g(z^{(L)}) = h_{\Theta}(x) \end{aligned} \quad (58)$$

The cost function (unregularised) is then defined as:

$$J_{\Theta}(x) = \frac{1}{m} \sum_{t=1}^m \sum_{k=1}^K [y_k^{(i)} \log(h_{\Theta}(x^{(i)})_k) + (1 - y_k^{(i)}) \log(1 - h_{\Theta}(x^{(i)})_k)] \quad (59)$$

And regularised is:

$$J_{\Theta}(x) = \frac{1}{m} \sum_{t=1}^m \sum_{k=1}^K [y_k^{(i)} \log(h_{\Theta}(x^{(i)})_k) + (1 - y_k^{(i)}) \log(1 - h_{\Theta}(x^{(i)})_k)] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{j,i}^{(l)})^2 \quad (60)$$

1.4.5 Backpropagation

The aim of back propagation is to calculate the gradient $\frac{\partial}{\partial \Theta_{i,j}^{(l)}} J(\Theta)$, which can then be used in an update rule such as gradient descent.

Back propagation is carried out element-wise through the training data, typically using a for-loop:

$$\begin{aligned} & \text{for } t \text{ in range}(m): \\ & \dots \\ & \dots \end{aligned} \quad (61)$$

Calculating δ

We calculate δ for layers $L, (L-1), (L-2), \dots, 2$. Note that we do not calculate δ^1 i.e. δ for the input layer.

For the output later:

$$\delta^{(L)} = a^{(L)} - y \quad (62)$$

For the hidden layers:

$$\delta^{(l)} = ((\Theta^{(l)})^T \delta^{(l+1)}) \cdot g'(z^{(l)}) \quad (63)$$

Where:

$$g'(z^{(l)}) = g(z^{(l)}) \cdot (1 - g(z^{(l)})) = a^{(l)} \cdot (1 - a^{(l)}) \quad (64)$$

And, as before (although written a little differently):

$$z^{(l)} = (\Theta^{(l-1)})^T a^{(l-1)} \quad (65)$$

And $\cdot *$ is the element-wise multiplication.

Note that δ^l is a 2-D vector with the same number of units as layer l :

$$[\delta^l] = \left[\begin{array}{c} \uparrow \\ s_l + 1 \\ \downarrow \end{array} \right] \quad (66)$$

Element-wise, this is implemented as:

for t in range(m) :

$$\begin{aligned} (\delta^{(L)})^{(t)} &= (a^{(L)})^{(t)} - y^{(t)} \\ (\delta^{(l)})^{(t)} &= ((\Theta^{(l)})^T \delta^{(l+1)}) \cdot * (a^{(l)})^{(t)} \cdot * (1 - (a^{(l)})^{(t)}) \end{aligned} \quad (67)$$

Where:

$$[(\delta^{(l)})^{(t)} = ((\Theta^{(l)})^T \delta^{(l+1)})] = \quad (68)$$

Calculating Δ

Note that $\Delta^{(l)}$ has the same dimensions as $\Theta^{(l)}$:

$$[\delta^{(l)}] = [\Theta^{(l)}] = \left[\begin{array}{ccc} & \leftarrow & s_l + 1 & \rightarrow \\ \uparrow & & & \\ s_{l+1} & & & \\ \downarrow & & & \end{array} \right] \quad (69)$$